

Voice Models: Claude vs ChatGPT vs Open Source

For the HR26 Agentic Voice AI problem. Facts as of 8 Oct 2026. ⚠ marks things I could not confirm from a primary source.

0. TL;DR

- **Claude voice mode** is a turn-based pipeline (listen → think → speak) with a Claude model as the brain. It is an app feature. I found no Claude voice API. ⚠
- **ChatGPT voice mode** now runs on **GPT-Live-1** (paid) / **GPT-Live-1 mini** (free): a full-duplex voice model (listens and talks at the same time) that hands hard reasoning to a backend model.
- **Neither consumer app can be your project's engine.** OpenAI exposes the same tech via API (GPT-Live-1, gpt-realtime-2.1). Anthropic does not expose Claude's voice stack.
- **Your constraint is zero API budget**, so the pick is an **open-source cascade**: Pipecat + faster-whisper (or Moonshine) + Ollama + Kokoro. Open speech-to-speech models (Moshi, PersonaPlex) are cool but the wrong risk for a same-day build.

1. What each one actually uses

1.1 Claude voice mode

Layer	What it is	Confidence
Architecture	Turn-based: Claude listens, pauses to process, then responds (Anthropic told Engadget this)	Reported
Brain (LLM)	Haiku / Sonnet / Opus; paid users pick the model, voice starts on the last model used in text. Fable not supported	Reported (23 Jul 2026 update)
Speech-to-text	Not disclosed	⚠
Text-to-speech	ElevenLabs technology, inferred from Anthropic's subprocessor list (found by Simon Willison). Not announced by Anthropic	⚠
Tools	Gmail, Google Calendar, Docs, Slack via connectors	Reported
Languages	English at launch, expanded (18 new languages in May 2026 per one source)	Reported
Camera/vision	Not as of May 2026	Reported

1.2 ChatGPT voice mode

Layer	What it is	Confidence
Model (app)	GPT-Live-1 (Go/Plus/Pro), GPT-Live-1 mini (Free). Rolled out 8 Jul 2026	Confirmed
Architecture	Full-duplex: speaks, hears and listens simultaneously; interruptions handled natively	Confirmed
Reasoning	Delegates search, reasoning, long tasks to a backend frontier model (GPT-5.5 at launch) while the conversation continues	Reported
Old mode	Advanced Voice Mode on a GPT-4o-era model (can still be selected in settings)	Reported
Limits	No screen/video sharing in GPT-Live at launch; some memory bugs reported	Reported

1.3 OpenAI's developer versions (what you could actually build on)

Model	Type	Price
GPT-Live-1 (API)	Full-duplex voice layer, can delegate to your own models/agents, telephony support	\$0.05/min front end, backend billed separately
gpt-realtime-2.1 / 2.1-mini	Speech-to-speech reasoning models (released 6 Jul 2026)	See OpenAI pricing page
gpt-realtime-2	Speech-to-speech, "GPT-5-class reasoning", 128K context, tool calls	\$32 / 1M audio tokens in, \$64 / 1M out
gpt-realtime-whisper	Streaming speech-to-text	\$0.017/min
gpt-realtime-translate	Live speech translation	\$0.034/min

2. Comparison

Dimension	Claude voice mode	ChatGPT voice (GPT-Live)
Conversation feel	Turn-based, noticeable pause	Full-duplex, natural interruptions
Reasoning ceiling	High (Opus/Sonnet selectable)	High (delegates to GPT-5.5-class backend)
Tool use by voice	Connectors (Gmail, Calendar, Slack, Docs)	Delegated agentic tasks; API supports tools
Usable in your project?	No (no voice API found) ⚠️	Only via paid API , not the app
Cost to you	Subscription + usage limits	API: paid per minute/token
Open / self-hostable	No	No
Vision in voice	No	No (screen share missing at launch)

Pros / cons

Claude

- ☒ Strong at careful, instruction-following reasoning; model choice in voice; good connector set
- ☒ Claude works well as the *LLM brain* inside someone else's voice pipeline
- ☒ Turn-based latency, no interruption handling like GPT-Live
- ☒ No public voice stack to build on; STT/TTS vendors not disclosed
- ☒ Voice turns burn your normal message quota

ChatGPT / OpenAI

- ☒ Best conversational naturalness right now (full-duplex)
- ☒ Same tech is available via API (GPT-Live-1, gpt-realtime-2.1) with tools, telephony, tone control
- ☒ Whisper-based streaming STT cheap at \$0.017/min
- ☒ Costs money per minute; you have no API budget
- ☒ Closed weights, cloud-only, data leaves the machine
- ☒ GPT-Live API is new, expect rough edges

3. Which is better for our project?

Assumption: the problem is "build an agentic voice AI": user speaks, agent plans, calls tools, speaks back. I don't have the official problem statement, so check it for any mandated models or platform rules.

Verdict by constraint:

If...	Pick	Why
You must choose between the two vendors	OpenAI	It is the only one with a developer voice API. Claude voice is app-only
You have API credits	GPT-Live-1 API or gpt-realtime-2.1	Best demo feel, least plumbing
You have no API budget (your stated setup)	Open-source cascade	Free, runs locally, you control tool calling
You want Claude's reasoning	Use Claude as the LLM in a cascade	Needs an API key, so only if you have credits

Why not the consumer apps: a hackathon judges *your* system. Driving the ChatGPT or Claude app by voice is a demo of their product, not yours.

4. Open-source options (free)

4.1 Pipeline pieces (cascade: STT → LLM → TTS)

Job	Model	Why use it	License	Link
Orchestration	Pipecat	Frame-based pipeline, built-in VAD, barge-in, any STT/LLM/TTS	BSD-2	https://github.com/pipecat-ai/pipecat
Orchestration	LiveKit Agents	WebRTC + telephony built in, open-weight turn-detection model	Apache-2.0	https://github.com/livekit/agents
STT	faster-whisper (large-v3-turbo)	Whisper accuracy, ~4x faster, 99 languages	MIT	https://github.com/SYSTRAN/faster-whisper
STT	Moonshine	Built for streaming; claimed sub-200ms on edge/CPU, English	MIT	https://github.com/moonshine-ai/moonshine
STT	Parakeet TDT v3	Very fast, strong English	CC-BY-4.0	https://huggingface.co/nvidia/parakeet-tdt-0.6b-v3
VAD	Silero VAD	Tiny, reliable speech-start/stop detection	MIT	https://github.com/snakers4/silero-vad
LLM	Ollama (Llama 3.1 8B / Phi-4-mini)	You already run it; tool calling for the "agentic" part	MIT	https://github.com/ollama/ollama
TTS	Kokoro-82M	Tiny, fast on CPU/GPU, good quality	Apache-2.0	https://github.com/hexgrad/kokoro
TTS	Chatterbox / Chatterbox-Turbo	Voice cloning, strong for voice agents	MIT	https://github.com/resemble-ai/chatterbox
TTS	Orpheus TTS	Emotion tags, streaming ~200ms claimed	Apache-2.0	https://github.com/canopyai/Orpheus-TTS

4.2 End-to-end speech-to-speech (no cascade)

Model	Why it's interesting	Why not for today	Link
Moshi (Kyutai)	First real-time full-duplex spoken LLM, ~200ms in practice	Small LLM brain, no native tool calling that I found ⚠️	https://github.com/kyutai-labs/moshi
PersonaPlex (NVIDIA, built on Moshi)	7B full-duplex, role + voice control via prompt, commercial-friendly	English-only; tool integration only "planned"; needs a real GPU	https://github.com/NVIDIA/PersonaPlex
Qwen3-Omni	Multilingual, function calling	80GB-class hardware for best results	https://github.com/QwenLM/Qwen3-Omni

4.3 "Is open source better than Claude/ChatGPT voice?"

- **Conversation naturalness:** No. GPT-Live is ahead of anything you can run locally today.
- **For this project:** Yes. It's free, private, scriptable, and tool calling is under your control. That matters more for an *agentic* build than a perfect voice.
- **Cost:** Cascade = \$0 marginal. OpenAI = per-minute.

5. Cheat sheet

- Claude voice = turn-based, app-only, TTS likely ElevenLabs ⚠️
- ChatGPT voice = GPT-Live-1, full-duplex, API exists at \$0.05/min + backend
- Can't build on either app; OpenAI is the only one with an API
- No budget → Pipecat + faster-whisper + Ollama + Kokoro
- Moshi / PersonaPlex: watch, don't bet today on them

Sources

- The Decoder: Claude uses ElevenLabs tech (<https://the-decoder.com/anthropics-claude-uses-elevenlabs-technology-for-speech-features-rather-than-an-in-house-model/>)
- Claude voice on Opus/Sonnet, 23 Jul 2026 (<https://letsdatascience.com/news/anthropic-extends-claude-voice-mode-to-opus-and-sonnet-16faf79b>)
- TechCrunch: Claude voice mode launch (<https://techcrunch.com/2025/05/27/anthropic-launches-a-voice-mode-for-claude>)
- Neowin: ChatGPT Voice on GPT-Live (<https://www.neowin.net/news/openais-next-generation-chatgpt-voice-will-make-advanced-voice-mode-look-outdated/>)
- OpenAI: GPT-Live-1 in the API (<https://openai.com/index/introducing-gpt-live-1-in-the-api/>)
- The Decoder: GPT-Realtime-2 (<https://the-decoder.com/openais-new-voice-model-brings-gpt-5-level-reasoning-to-real-time-conversations/>)
- Neowin: GPT-Realtime pricing (<https://www.neowin.net/amp/openai-unveils-trio-of-realtime-audio-models-to-power-next-gen-voice-agents/>)
- DataNorth: gpt-realtime-2.1 (<https://datanorth.ai/news/openai-releases-gpt-realtime-2-1-voice-models>)
- Open Source For You: PersonaPlex (<https://www.opensourceforu.com/2026/08/nvidia-open-sources-personaplex-built-on-moshi/>)
- Moshi paper (<https://huggingface.co/papers/2410.00037>)
- Codesota open-source TTS ranking (<https://codesota.com/speech/best-open-source>)
- Northflank open-source STT guide (<https://northflank.com/blog/best-open-source-speech-to-text-stt-model-in-2026-benchmarks>)
- Pipecat vs LiveKit (<https://www.dograh.com/feeds/blog/pipecat-livekit-agents>)