

HR26: Agentic Voice AI Problem Statements

Problem statements with Agnes models and API access

HR26 AI Track | Agnes AI India Hackathon | 8 October 2026

Participant Note and Common API Access

Note for Participants: The AI Track includes two categories: **Voice AI** and **Agentic AI**. Teams may choose **any one of the problem statements below** and are free to implement voice interfaces, agentic workflows, or a combination of both.

The API details below apply to all three problems; each problem then adds its own model mapping.

Item	Detail
Platform	https://platform.agnes-ai.com
Base URL	https://apihub.agnes-ai.com/v1
Text	POST /v1/chat/completions or POST /v1/responses (Anthropic-compatible POST /v1/messages also available)
Images	POST /v1/images/generations
Video	POST /v1/videos, retrieved at https://apihub.agnes-ai.com/agnesapi using video_id and model_name=agnes-video-2.5
Free-plan RPM	Text 10 · Image 1K 10 · Image 2K 5 · Image 3K/4K 1 · Video 1
Rules	Keep API keys on the server. Queue requests and use backoff on rate limits. Video is asynchronous, so show task status and retrieve the finished output before reporting success. Label saved outputs used as a demo fallback.

HR26-AI-01: Learning Experiences

The Problem

Educators rarely think about teaching while sitting at a desk. Ideas arrive in staff rooms, on commutes and between classes, often spoken out loud, half-finished and subject to change. What they have to work with is fragmented: lecture notes, PDFs, textbooks, presentations, reference links and classroom requirements, with very little time to turn any of it into something that works for their learners.

Learners differ in prior knowledge, language, pace and the way they engage with a subject. Some read comfortably, some learn better by listening and speaking, and some are not yet confident in the language of instruction. What works for one classroom may not work for another, and requirements can

change while preparation is underway, as when a class falls behind, an exam moves or a new student joins.

The challenge is to rethink how educators and learners can turn available knowledge, resources and teaching goals into meaningful learning experiences, through the way people naturally talk about learning.

The solution should enable an educator to go from scattered resources and rough goals to a learning experience they trust and are ready to use, while faithfully reflecting their own material. It should work for different prior knowledge, languages, paces and ways of engaging, and adapt when requirements change without starting over. Educators should remain in control, be able to see where understanding is growing or not, distinguish what comes from their material, and operate within realistic classroom constraints of time, connectivity and cost.

Minimum Objectives

- Adapt learning content based on learner performance, prior knowledge and preferred learning style.
- Predict future learning gaps and proactively redesign the learner's path before those gaps emerge.
- Resolve conflicting or outdated information across multiple sources and autonomously determine what knowledge should be trusted.

Open Creativity

The Minimum Objectives are the minimum that must be achieved by your solution.

Beyond these requirements, you are encouraged to think beyond the obvious, explore unconventional approaches, identify additional challenges within the problem and introduce your own innovative capabilities.

The more creatively you solve the problem, the more you think out of the box, and the further you push the solution beyond the stated objectives, the more points you can earn.

Models & API Access

Model	Identifier	Role in this problem
Agnes 3.0 Flash (primary)	agnes-3.0-flash	The 512K context holds lecture notes, PDFs and textbooks together. Use it for source comparison, trust ranking, learner-path planning and multilingual content, with tool calling for the agent loop.
Agnes Image 2.5 Flash (optional)	agnes-image-2.5-flash	Diagrams and visual aids for learners who prefer visuals.

- **Endpoints:** `/v1/chat/completions` for the main agent. `/v1/images/generations` only if you add visuals.
- **Pricing (standard):** Text is \$0.05 per 1M input tokens (\$0.005 cached) and \$0.15 per 1M output tokens. Images are \$0.010 (1K) to \$0.024 (4K) each.
- **Free access:** Both models are free on the platform, within the RPM limits above.

- **Note:** Agnes 3.0 Flash takes text and image-URL input only. Speech-to-text and text-to-speech for a voice interface would need a separate component.

HR26-AI-02: Marketing Campaigns

The Problem

A café owner in Bengaluru knows exactly what she wants people to hear: a weekend offer, in a tone that sounds like her, in the languages her customers use. But she is behind the counter, her hands are busy and her ideas change as she sees what takes shape. She has no designer, no marketing team and no time to write a brief.

Small businesses and brands often have a clear idea of what they want to communicate but lack the time, resources or expertise to turn it into an effective campaign. A single campaign may need to carry the same message to different audiences, in different languages, across different channels, while staying recognizable and consistent. Requirements keep changing as owners see and react to what is being created, and the one thing that cannot go wrong is what customers are told to expect.

The challenge is to rethink how businesses can turn their ideas and objectives into effective marketing experiences that remain coherent as the campaign evolves.

The solution should enable a business owner to go from a half-formed idea to a campaign they are proud to put their name on, with very little effort. The message, offer and brand should remain consistent across audiences, languages and channels, while ensuring that nothing customers are told to expect is wrong. Changing requirements should be absorbed without losing what was already approved. The owner should know what exists, what changed and what remains pending, while the result feels local and authentic rather than generic or translated. The owner should remain in control, with realistic time and cost requirements for a small business.

Minimum Objectives

- Generate a coherent campaign adapted to different audiences, languages and channels.
- Predict campaign performance before launch and autonomously optimize the campaign to maximize its effectiveness.
- Preserve the exact business intent while autonomously adapting messaging to different cultural and linguistic contexts.

Open Creativity

The Minimum Objectives are the minimum that must be achieved by your solution.

Beyond these requirements, you are encouraged to think beyond the obvious, explore unconventional approaches, identify additional challenges within the problem and introduce your own innovative capabilities.

The more creatively you solve the problem, the more you think out of the box, and the further you push the solution beyond the stated objectives, the more points you can earn.

Models & API Access

Model	Identifier	Role in this problem
Agnes 3.0 Flash (primary)	agnes-3.0-flash	Campaign planning, copy in several languages and tones, a locked "offer facts" record checked against every output, performance prediction and optimization.
Agnes Image 2.5 Flash	agnes-image-2.5-flash	Posters and social creatives. Supports editing and multi-image composition, so brand assets can be reused across variants.
Agnes Video 2.5 or Video 2.5 Flash (optional)	agnes-video-2.5	Short promo clips (4–12 s). Flash is free for a limited time and 720P only.

- **Endpoints:** /v1/chat/completions, /v1/images/generations, and /v1/videos for clips.
- **Pricing (standard):**
 - Text: \$0.05 per 1M input tokens and \$0.15 per 1M output tokens.
 - Images: \$0.010 (1K), \$0.018 (2K), \$0.021 (3K), \$0.024 (4K). Reference images beyond the first three add \$0.003 each.
 - Video 2.5: \$0.025 per second at 720P, \$0.040 at 1080P or 1K, and \$0.055 at 2K. Reference images beyond the first five add \$0.005 each.
- **Free access:** Text, Image and Video 2.5 Flash are free on the platform. Standard Video 2.5 is paid.
- **Note:** At 1 RPM, video should be queued. Keep a clearly labeled saved clip as a demo fallback.

HR26-AI-03: Video Content Creation

The Problem

A teacher records an hour-long lesson. A founder films a rambling product walkthrough. A community organizer interviews villagers in three languages. In each case the raw footage contains something valuable, but buried among repetition, mistakes, tangents and moments meant for different audiences. Turning it into something people will actually watch takes time, effort and creative expertise that most creators do not have.

The spoken word is often the heart of that footage: what was said, what was meant, what mattered and what the speaker would never want presented out of context. Understanding what a recording is trying to communicate, and shaping it for the people who will watch it, is usually a slow, manual process, and it starts again whenever the creator's goals or audience change.

The challenge is to rethink how raw video can be transformed into valuable content that serves the creator's intent and the needs of its intended audience.

The solution should transform raw footage into content that faithfully reflects the creator's intent, context and important information, without misrepresentation. One recording should be usable for different audiences and purposes without being rebuilt each time. Creators should be able to steer and revise the outcome as their goals change, while handling real-world footage containing mistakes,

repetition, tangents, multiple speakers, mixed languages and noise. Creators should also be able to verify what was kept, what was cut and why, while keeping time and cost practical for individuals and small organizations.

Minimum Objectives

- Generate audience-specific content from raw footage while preserving the creator's intent and important information.
- Understand implicit meaning and preserve its context while automatically restructuring and heavily compressing the footage.
- Guarantee that automatically generated edits never introduce misleading context or change the intended meaning.

Open Creativity

The Minimum Objectives are the minimum that must be achieved by your solution.

Beyond these requirements, you are encouraged to think beyond the obvious, explore unconventional approaches, identify additional challenges within the problem and introduce your own innovative capabilities.

The more creatively you solve the problem, the more you think out of the box, and the further you push the solution beyond the stated objectives, the more points you can earn.

Models & API Access

Model	Identifier	Role in this problem
Agnes 3.0 Flash (primary)	agnes-3.0-flash	Reads the transcript to find intent, implicit meaning, repetition and tangents. Builds an edit decision list per audience, with a kept/cut/why log and a meaning-preservation check on every cut.
Agnes Image 2.5 Flash (optional)	agnes-image-2.5-flash	Thumbnails, title cards and cover images.
Agnes Video 2.5 or Video 2.5 Flash (optional)	agnes-video-2.5	Generates short intro/outro or b-roll clips (4-12 s). It is not an editor for the creator's own footage.

- **Endpoints:** /v1/chat/completions for analysis, /v1/images/generations for covers, /v1/videos for generated clips.
- **Pricing (standard):** Text and image rates are the same as above. Video is billed per second for both output and input video duration, so feeding long raw footage as a reference gets expensive.
- **Important limit:** Agnes 3.0 Flash accepts text and image-URL input only, not audio or video. You will need a separate speech-to-text step to produce a timestamped transcript. The actual cutting and compressing should be done with a deterministic tool such as ffmpeg, using the model's cut list. Cutting the creator's real footage, and not regenerating it, is what makes the "never misrepresent" guarantee achievable.

Official Resources

Resource	Link
Website	agnes.ai
API platform	https://platform.agnes-ai.com/
Agnes 3.0 Flash documentation	https://wiki.agnes-ai.com/en/docs/agnes-30-flash
Agnes Image 2.5 Flash documentation	https://wiki.agnes-ai.com/en/docs/agnes-image-25-flash
Agnes Video 2.5 documentation	https://wiki.agnes-ai.com/en/docs/agnes-video-25
Current model pricing	https://wiki.agnes-ai.com/en/docs/pricing
Access plans and RPM limits	https://wiki.agnes-ai.com/en/docs/tokenplan
Agnes Code	https://agnes-ai.com/agnescode
Model repository	https://github.com/AgnesAI-Labs/AgnesAI-Models
Agnes Harness	https://github.com/AgnesAI-Labs/agnes-harness
Community Discord	https://discord.gg/agnesai
X	https://x.com/agnesai_sapiens
LinkedIn	https://www.linkedin.com/company/agnesaisapiens/

Media and partnership inquiries: daniel@sapiens-ai.io

Prices and RPM limits are as of the Agnes brief dated 2 October 2026, so check the platform before the event.