

Problem

Small business owners know what they want to say but have no time, designer or marketer. Campaigns have to stay consistent across languages and channels, and the offer details can never be wrong.

Target user

Priya, who runs a café in Indiranagar. She's behind the counter with busy hands. Her customers speak Kannada, Hindi and English, and she posts on Instagram and WhatsApp.

Goals → minimum objectives

Objective	How we meet it
Coherent campaign across audiences, languages, channels	Planner matrix + shared brand voice + Offer Facts
Predict performance and autonomously optimize	Persona panel scoring + optimizer loop
Preserve exact intent while adapting culturally	Native-language generation + deterministic validator + back-translation drift check

User stories

1. Priya speaks her idea and gets a full campaign in minutes
2. She sees and approves the locked offer details before anything goes out
3. She says "make it Sunday only" and only the affected assets change
4. She can see what's approved, what changed and what's pending
5. She sees which variant will likely do best and why
6. She's told about any asset that got blocked for a wrong price or date

Features

P0 (must ship)

- Voice input (EN + KN/HI) → brief
- Offer Facts lock + approval
- Copy in 3 languages × 3 channels
- Posters (Agnes image + text overlay)
- Deterministic fact validator
- Persona scoring + one optimize loop
- Voice change → selective regeneration
- Status board + change log

P1 (if time)

- 8 s promo reel via Agnes Video Flash

- Voice read-back summary (TTS)
- Brand kit memory (logo, colours, tone) reused across campaigns

P2 (pitch only)

- Real post scheduling, live engagement feedback for real performance learning, festival calendar suggestions

Out of scope

Auto-posting to real accounts, paid ad buying, analytics integrations.

Success metrics (demo)

- Voice to full campaign in under 3 min
- 0 fact mismatches reach "Approved"
- A change touches only the dependent assets
- The optimized variant visibly beats the original's persona score

36-hour plan

- **0–4 h:** API keys, STT test, skeleton repo, Offer Facts schema
- **4–14 h:** Brief, Planner and Copy agents + validator
- **14–22 h:** Posters + overlay, persona panel, optimizer
- **22–28 h:** Change handling + board UI
- **28–32 h:** Video (P1), polish, cache demo outputs
- **32–36 h:** Rehearse the demo, pitch deck, fallback recordings

Risks

- Indic STT accuracy → show the transcript and allow text edits
- Rate limits → queue + cache + saved fallbacks
- Persona scores seen as fake → frame them honestly as a pre-launch proxy with visible reasoning